

The Build Fellowship

BUILDFELLOWSHIP.COM



Open Avenues

Weekly Updates

- Please provide a quick update on either:
 - Something you did/saw this week that you thought was interesting
 - What you're looking forward to about this week's workshop

(Reminder - please have your cameras on if possible)



The Build Fellowship

Workshop 5

Model Training

Approaches

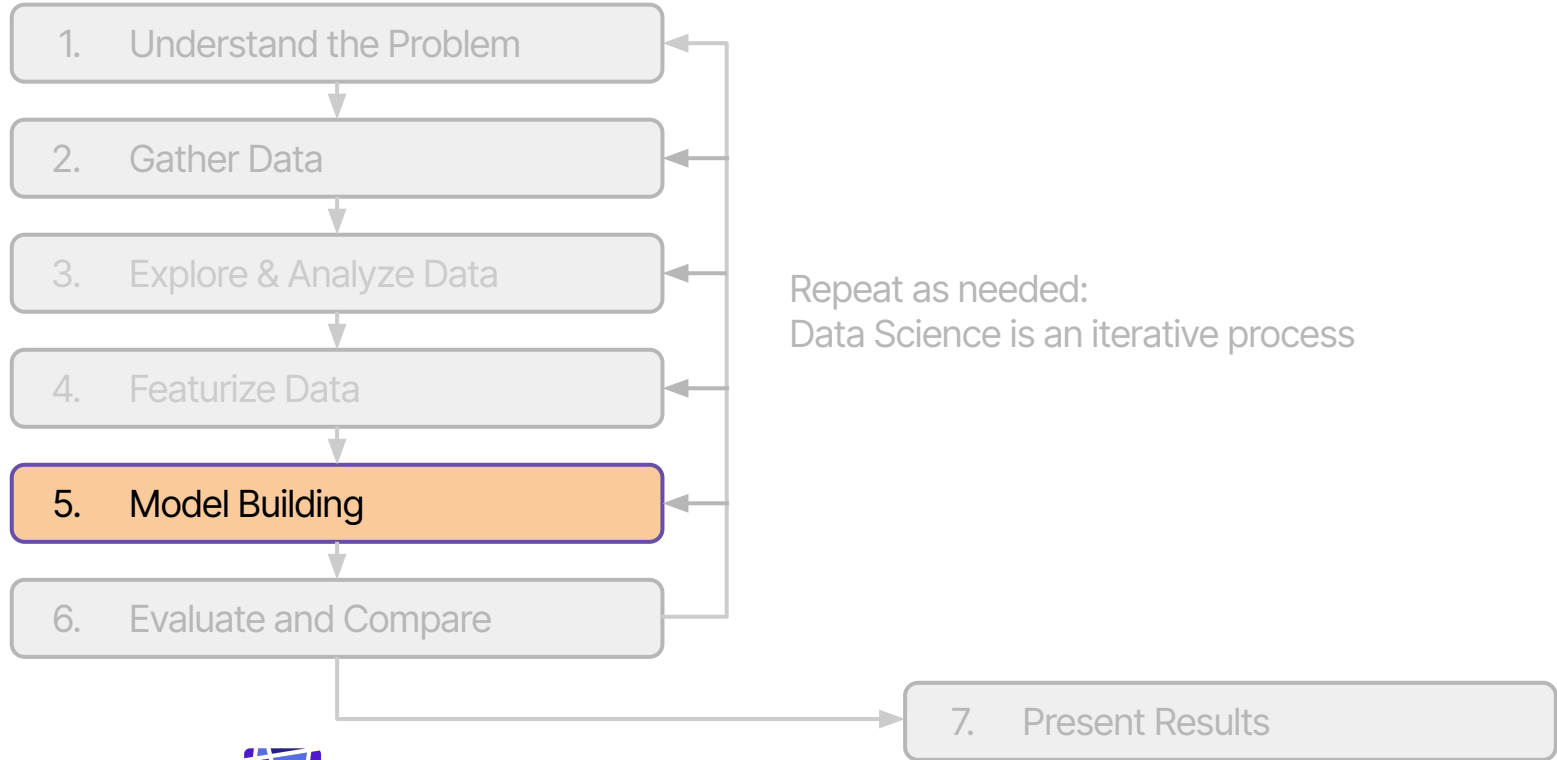
Recap



Sessions Overview

- Workshop 1 – Project Introduction & Setup
- Workshop 2 – Genomic Data (A2 Assignment)
- Workshop 3 – Data Analysis & Visualization (A3 Assignment)
- Workshop 4 – Featurization & Baseline Modeling (A4 Assignment)
- **Workshop 5 – Model Training Approaches (Final Assignment Set)**
- Workshop 6 – Model Tuning
- Workshop 7 – Performance Evaluation (Final Assignment Code/Testing Due)
- Workshop 8 – Results Presentation & Wrap up (Final Presentation Due)

The Data Science Process



From Features to Models

- Last week we explored a number of featurization methods
- 3 possible feature spaces to leverage for modeling
 - Opportunity to expand further
- Assignment 4 - already built a simple model
- What comes next?
 1. Review modeling options
 2. Relate models to feature options
 3. Determine a robust framework for training models

Model Types



Linear Models

- Binary S vs R targets
- Good starting point


Examples:

- Regression
 - Linear
 - Logistic
- SVM

Assumptions:

- Linear relationships
- Is this expected from our data?
 - Presence/Absence
 - Kmers
 - Sequences

Continuous Linear Model



Presence Absence Binary Model



Logistic Regression

- Simple Linear regression won't work for us
- Target is binary & Features = binary or counts


Overview:

- Logistic regression = GLM
- Link function
- Mapping from continuous to bounded
- Sigmoid function

Python:

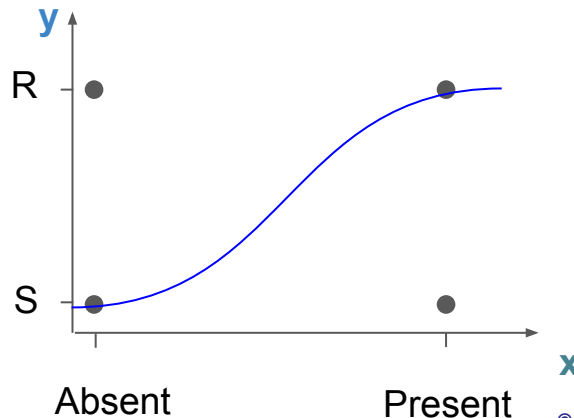
- `Sklearn.linear_model`
- [Documentation](#)

Logistic Model



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Presence Absence Binary Model



$y =$
"predicted
probability"

Tree Based Models

- Decision Trees
- Random Forest (Ensemble)

Decision Trees:

- Non-linear modeling approach
- Based off stacking binary decisions
- Natural fit for our presence/absence features
- Decision trees rarely used alone


Random Forest:

- Ensemble = combine multiple models
- Leverages Bagging (Bootstrap aggregating)

Python:

- Sklearn.ensemble
- [Documentation](#)

Decision Tree



Process:

1. Decide on which feature to split (gain)
2. Determine a splitting threshold
3. Continue splitting until stop criteria (e.g. maximum depth)

Inference:

- New sample traverses tree
- Ends up at a single “leaf node”
- Take average value from leaf

Bagging = Repeatedly sample with replacement

Tree Based Models

- Decision Trees
- Random Forest (Ensemble)

Decision Trees:


- Non-linear modeling approach
- Based off stacking binary decisions
- Natural fit for our presence/absence features
- Decision trees rarely used alone

Random Forest:

- Ensemble = combine multiple models
- Leverages Bagging (Bootstrap aggregating)

Python:

- Sklearn.ensemble
- [Documentation](#)



Tree Based Models

- Decision Trees
- Random Forest (Ensemble)

Decision Trees:

- Non-linear modeling approach
- Based off stacking binary decisions
- Natural fit for our presence/absence features
- Decision trees rarely used alone


Random Forest:

- Ensemble = combine multiple models
- Leverages Bagging (Bootstrap aggregating)


Python:

- Sklearn.ensemble
- [Documentation](#)

Bagging = Repeatedly sample with replacement



Random Forest = Build New Tree per Bag



Tree Based Models

- Decision Trees
- Random Forest (Ensemble)

Decision Trees:

- Non-linear modeling approach
- Based off stacking binary decisions
- Natural fit for our presence/absence features
- Decision trees rarely used alone


Random Forest:

- Ensemble = combine multiple models
- Leverages Bagging (Bootstrap aggregating)


Python:

- Sklearn.ensemble
- [Documentation](#)


Bagging = Repeatedly sample with replacement



Random Forest = Build New Tree per Bag



Prediction = Each Sample passes through each Tree



Output = Average across trees

Gradient Boosting


- Ensemble method
- Can use any underlying model
 - Common to use decision trees

Overview:

- Iteratively fit models on residuals of the previous
- Learn from the mistakes of the last model
- Leveraging lots of simple models

Options:

- [Sklearn Gradient Boosting](#)
- [LightGBM](#)
- [XGBoost](#)



Gradient Boosting



- Ensemble method
- Can use any underlying model
 - Common to use decision trees

Overview:

- Iteratively fit models on residuals of the previous
- Learn from the mistakes of the last model
- Leveraging lots of simple models

Options:

- [Sklearn Gradient Boosting](#)
- [LightGBM](#)
- [XGBoost](#)



Gradient Boosting


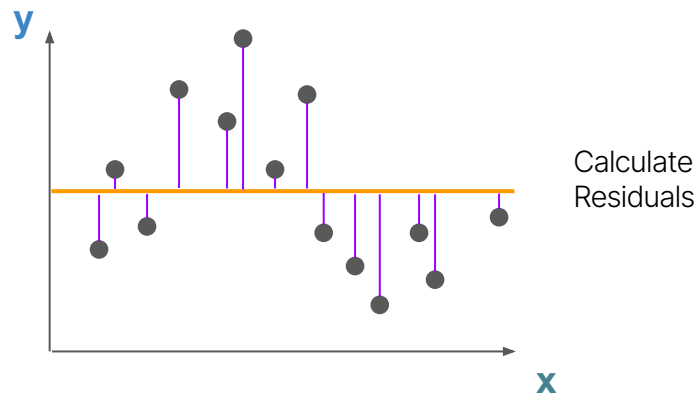
- Ensemble method
- Can use any underlying model
 - Common to use decision trees

Overview:

- Iteratively fit models on residuals of the previous
- Learn from the mistakes of the last model
- Leveraging lots of simple models

Options:

- [Sklearn Gradient Boosting](#)
- [LightGBM](#)
- [XGBoost](#)



Neural Networks

- Many different architectures
 - Convolutional Neural Network (CNN)
 - Recurrent Neural Network (RNN)
- Tabular & Sequence data (e.g. audio/image)


Overview:

- Very powerful and flexible
- Expensive to train & can overfit easily
- Network of simple nonlinear functions

Options:

- [Tensorflow](#)
- [Keras](#) (higher level interface)
- [Pytorch](#)

Single Neuron



$$y = x1 + x2$$

Neural Networks

- Many different architectures
 - Convolutional Neural Network (CNN)
 - Recurrent Neural Network (RNN)
- Tabular & Sequence data (e.g. audio/image)


Overview:

- Very powerful and flexible
- Expensive to train & can overfit easily
- Network of simple nonlinear functions

Options:

- [Tensorflow](#)
- [Keras](#) (higher level interface)
- [Pytorch](#)

Single Neuron



$$y = w1x1 + w2x2$$

Neural Networks

- Many different architectures
 - Convolutional Neural Network (CNN)
 - Recurrent Neural Network (RNN)
- Tabular & Sequence data (e.g. audio/image)


Overview:

- Very powerful and flexible
- Expensive to train & can overfit easily
- Network of simple nonlinear functions

Options:

- [Tensorflow](#)
- [Keras](#) (higher level interface)
- [Pytorch](#)

Single Neuron



$$y = w1x1 + w2x2 + b$$

Neural Networks

- Many different architectures
 - Convolutional Neural Network (CNN)
 - Recurrent Neural Network (RNN)
- Tabular & Sequence data (e.g. audio/image)


Overview:

- Very powerful and flexible
- Expensive to train & can overfit easily
- Network of simple nonlinear functions

Options:

- [Tensorflow](#)
- [Keras](#) (higher level interface)
- [Pytorch](#)

Single Neuron



$$y = f(w_1x_1 + w_2x_2 + b)$$

Neural Networks

- Many different architectures
 - Convolutional Neural Network (CNN)
 - Recurrent Neural Network (RNN)
- Tabular & Sequence data (e.g. audio/image)


Overview:

- Very powerful and flexible
- Expensive to train & can overfit easily
- Network of simple nonlinear functions

Options:


- [Tensorflow](#)
- [Keras](#) (higher level interface)
- [Pytorch](#)

Single Neuron



$$y = f(w_1x_1 + w_2x_2 + b)$$

Neural Network



$$y = f(f() + f() + \dots)$$

Neural Networks

- Many different architectures
 - Convolutional Neural Network (CNN)
 - Recurrent Neural Network (RNN)
- Tabular & Sequence data (e.g. audio/image)

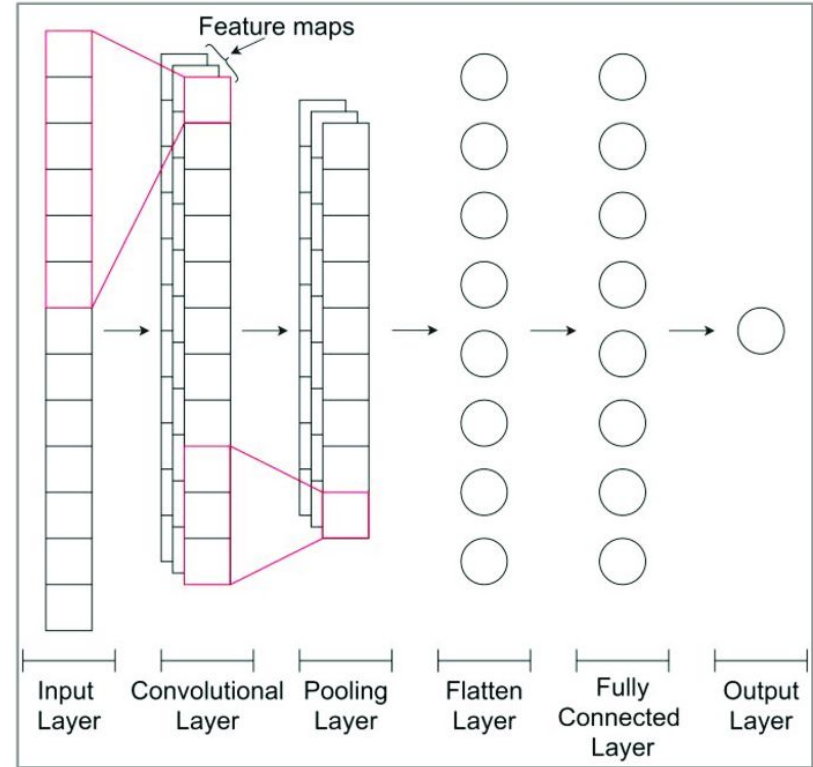
Overview:

- Very powerful and flexible
- Expensive to train & can overfit easily
- Network of simple nonlinear functions

Options:

- [Tensorflow](#)
- [Keras](#) (higher level interface)
- [Pytorch](#)

Convolutional (1D) Neural Network



[1]

QUIZ TIME !?

Which model would you try first and why?

- a) Convolutional neural network
- b) Decision tree
- c) Logistic regression
- d) Random forest
- e) Gradient boosted trees



Model Training



Data Splitting - Cross validation

Cross validation is a fundamental concept for training ML models:

- Training a single model is risky
- Wish to train many models to study uncertainty

Dataset



Standard splitting method: **K-fold**

Process:

1. Shuffle your data
2. Split the data into $1/K$ chunks
3. Iteratively split $1/K$ to validate and $(K-1)/K$ to train

Each samples appears in validate exactly once

Note: separate from having a completely held out test set

Data Splitting - Cross validation

Cross validation is a fundamental concept for training ML models:

- Training a single model is risky
- Wish to train many models to study uncertainty

Standard splitting method: **K-fold**

Process:

1. Shuffle your data
2. Split the data into $1/K$ chunks
3. Iteratively split $1/K$ to validate and $(K-1)/K$ to train

Each samples appears in validate exactly once

Note: separate from having a completely held out test set

Dataset



Shuffle ordering



Data Splitting - Cross validation

Cross validation is a fundamental concept for training ML models:

- Training a single model is risky
- Wish to train many models to study uncertainty

Standard splitting method: **K-fold**

Process:


1. Shuffle your data
2. Split the data into $1/K$ chunks
3. Iteratively split $1/K$ to validate and $(K-1)/K$ to train

Each samples appears in validate exactly once

Note: separate from having a completely held out test set

Say **K = 3**

Split the Data into 3 chunks



Data Splitting - Cross validation

Cross validation is a fundamental concept for training ML models:

- Training a single model is risky
- Wish to train many models to study uncertainty

Standard splitting method: **K-fold**

Process:


1. Shuffle your data
2. Split the data into $1/K$ chunks
3. Iteratively split $1/K$ to validate and $(K-1)/K$ to train

Each samples appears in validate exactly once

Note: separate from having a completely held out test set

Say **K = 3**

Split the Data into 3 chunks



Data Splitting - Cross validation

Cross validation is a fundamental concept for training ML models:

- Training a single model is risky
- Wish to train many models to study uncertainty

Standard splitting method: **K-fold**

Process:


1. Shuffle your data
2. Split the data into $1/K$ chunks
3. Iteratively split $1/K$ to validate and $(K-1)/K$ to train

Each samples appears in validate exactly once

Note: separate from having a completely held out test set

Say **K = 3**

Split the Data into 3 chunks



Repeat 3 times to make 3 datasets
(Orange box = Validate)

Final Project



Final Project - Overview

Challenge: **Build the best model you can to Predict Cefepime S vs R in Escherichia coli**

Requirements

- Use any feature sets
- Try at least two models
- Use CV
- Tune model to optimize performance
- Evaluate and discuss model performance

Future weeks will cover tuning and evaluation in more detail

Final Project - Evaluation

Challenge: **Build the best model you can to Predict Cefepime S vs R in Escherichia coli**

Evaluation

- Not being scored based on model performance
 - Should be able to get good (>90% accuracy) performance using simple models + features
1. Clear and well documented code
 2. Well specified and trained model
 3. Discussion and presentation of results

Submission:

- Results Notebook (can include .py file as desired)
- Summary slides (5-10 slide max, 5 minute run through)

Final Project - Evaluation

Challenge: **Build the best model you can to Predict Cefepime S vs R in Escherichia coli**

- <https://www.kaggle.com/t/e4838697a45d4d1dac9dfca2fe16d687>
- Join using the link above!
- This should be very lightweight and happy to answer any questions about using the page

Workshop 5

Model Building



References

[1]: https://www.researchgate.net/figure/One-dimensional-convolutional-neural-network-1D-CNN-architecture-for-the-timeseries_fig2_348502722

